HIGHLIGHTS:

- A novel method for mitigating bias in discriminative dataset by using an adversarial framework to learn an instance reweighing function is proposed

- FAIR comes in four variants differing in the weighting method

- FAIR is evaluated on four different real-world datasets and compared to the state-of-the-art techniques.

- The method can provide interpretable information about fairness of individual instances